Take another look at project ideas: https://fullstackdeeplearning.com/course/2022/project-showcase/ and https://fullstackdeeplearning.com/spring2021/projects/

(7p) **A.** Your project needs to be an AI based one, with online / cloud / parallelized deployment, accessible by clients remotely.

Check https://fullstackdeeplearning.com/course/2022/lab-7-web-deployment/ for deployed and automated GUIs.

- It would be nice to have cloud deployment. Use free accounts on either azure, google cloud, amazon, whatever your like. You must prove you can train on the cloud and serve your client's requests using models deployed on it. If not using cloud, the bare minimum requirement is to use Kubernetes orchestration or anything like prove scalability at least theoretically.

Check https://fullstackdeeplearning.com/spring2021/lecture-11/ for remote deploying ideas and strategies.

(1p) B. Experiment management

- Show that you have don experiment management at some level using TensorBoard or wand. Check the lab or course: https://fullstackdeeplearning.com/course/2022/lab-4-experiment-management/

(2p) C. Adopt a continual learning strategy and simulate that new data arrives sometimes.

Decide yourself what you do with the new data, either retraining everything or simply fine tune. Look again at https://fullstackdeeplearning.com/course/2022/lecture-6-continual-learning/

Please provide metrics on the differences between retraining vs fine-tuning when it comes to performance (time needed to retrain) and quality metrics such as accuracy, precision, recall, whatever you consider.

(2p) D. Documentation of your project with 10-15 slides that you must present along with the application, like the links containing project-showcases given above.

Note:

The total is 12p. The reason is that you can chose what you like to do. But presentation and base project is mandatory.